

7. SQL – Data Handling

7.1 The query language SQL

- Search predicates
- Arithmetic expressions and functions in predicates
- Different kinds of join
- Output layout

7.2 Advanced SQL

- Subselects and Correlated subqueries
- Quantified expressions, SOME, ANY
- Grouping and Aggregation
- Transitive closure

7.3 Update, Deletion, Insertion and bulk load*

Lit.: Melton / Simon, Understanding SQP 1999, chap. 2.5.7; Kemper / Eickler chap 4, SQL chapter in any book on DBS

(*) chap.6. Calculus Language: not discussed in class

SQL / DML: Overview

Query data

- Interactively
- Embedded in host language
important in applications, next chapter

Insert, update, delete data

© HS-2010

07-DBS-SQL-2

7.1 The Query Language SQL

SQL is relational complete

...but many additional query concepts compared to RA

Advanced search predicates on strings

e.g., find all cities starting with "Ber"

Arithmetics in expressions,

e.g., GNP / population for all countries

Grouping and predicates over sets

e.g., total GNP of EU countries

Recursion

© HS-2010

07-DBS-SQL-3

Simple SQL Search predicates

Defined like Boolean predicates in RA and Calculus
Some syntax extensions

- **<attribute> BETWEEN <value1> AND <value2>**
- **<attribute> IS [NOT] NULL**

```
SELECT *
FROM Country
WHERE population BETWEEN 50000000 AND 70000000
AND GNP IS NOT NULL
```

© HS-2010

07-DBS-SQL-4

Simple SQL Search expressions

Simplified OR

```
SELECT Name
FROM Country
WHERE C_ID IN ('BRD', 'D', 'DDR')
```

Equivalent to

<attr> = <val1> OR <attr> = <val2> OR ...

More general case of a **table constant**:

```
SELECT Name
FROM Country
WHERE (C_ID, capital)
IN (('BRD', 'Bonn'),
   ('D', 'Bonn'),
   ('DDR', 'Berlin (Ost)'))
```

© HS-2010

07-DBS-SQL-5

Boolean logic ...

.. is crucial:

```
SELECT Country.Name
FROM Country
WHERE NOT (capital != 'Berlin' OR
capital != 'Wien')
```

Result?

Row predicates with conjunctive primitives

a = <value> AND a = <differentValue> will never have a non-zero result.

© HS-2010

07-DBS-SQL-6

Remember

All logical differences are big differences
(Wittgenstein)

Corollar:

All logic mistakes are big miscakes

not so big

3-valued logic

NULL values

- comparison result may be "unknown" if an argument is NULL
 - "unknown" result tuples in any result set
- comparison only by IS [NOT] NULL

^	t	u	f
t	t	u	f
u	u	u	f
f	f	f	f

or, not
?

```
(select Country
from Economy
where gdp >= 0
) intersect
(select Country
from Economy
where gdp IS NULL)
```

Result set : Ø

Arithmetic functions in search predicates and more

Extensions of RA - Arithmetic expressions

May occur in simple predicates and target list

Basically arithmetic expressions (over attributes and values) allowed when attribute names allowed.

```
SELECT c.name,
       e.gdp/c.population AS "GDP ($ per Person"
FROM Country c join Economy e
ON c.code = e.country
WHERE (gdp*1000000/population) < 500
```

Result is NULL if any involved attribute is NULL

Simple SQL String search expressions

String expressions

Simple form of regular expressions: LIKE

LIKE patterns:

% : any sequence of characters
_ : exactly one character

```
SELECT Country.name , capital
FROM Country
WHERE capital LIKE '%co%'

NAME                CAPITAL
-----
Russia              Moscow
...
```

Regular expressions defined in SQL99*

<string> SIMILAR TO <pattern>
REGEXP_Like (<attr>, <pattern>) (Oracle)

Simple SQL Functions in search expressions

Built in functions

- Expressions may contain functions
- Many arithmetical and string built-in functions
- User defined functions on user defined types (see below)

```
SELECT Country.name , capital
FROM Country
WHERE SOUNDEX(capital) = SOUNDEX
('Monakko')
```

NAME	CAPITAL
Nicaragua	Managua
Monaco	Monaco

Simple SQL Search expressions

More expressions

Arithmetic expressions

example

Constant	7.5, 3.
[qualifier.]columnname	T.format, fname
arithExpr op arithExp	3 + 4, price - 2*discount
function (aexpr)	Sqrt(xCoord*xCoord + yCoord * yCoord)

Character and Dateexpressions

cexpr '+' cexpr	'The winner is:' + fname
function (cexpr)	SOUNDEX ('Meier')
	UPPER (fname), LOWER(...)
	TRIM (TRAILING ' ' FROM ' Hello ')
	SUBSTRING(fname FROM 0 FOR 4)
Date value express.	SYSDATE - date_of_birth

Date-Functions

SQL **date model** based on timestamps with/out time zone,
types: **timestamp**, **date**, **time** (of day), **interval**.

```
SELECT bike_ID, year_bought
FROM Bikes
WHERE MONTHS_BETWEEN(SYSDATE, year_bought) > 24 ;

SELECT Bike_ID
FROM Bikes
WHERE TO_DATE('1.1.2009') > year_bought
```

Problem: **compatibility**, e.g. functions on time values
General issue:

Casting may result in not a well defined value
e.g. a time interval of one year and five month to seconds
(how many month with 30 | 31 | 28 days?)

Naming

User defined type components
examples: Postgres

```
CREATE TYPE Coord AS (
  longitude NUMERIC,
  latitude NUMERIC)

CREATE TABLE City AS (
  name VARCHAR(..),
  coordinates Coord,
  ....
```

How to access components?

```
SELECT name FROM City
WHERE Coordinates.latitude = 37.5
```

Does not work:
looks like path expression

```
SELECT name FROM City
WHERE (Coordinates).latitude = 37.5
```

Array types much more involved! See manual.

Constructing Joins

Recipe for more involved join

Example: Find countries having cities with a population of 5,000,000 and more; list also city names and Province name and population.

1. Construct a "wide table" by joining all "relevant" tables.
2. Apply selection and projection to this table.

```
J-Tab(...Country.code,
      Country.name, ...
      .. Province.country,
      .. Province.name,
      ....
      City.country, -- code
      City.province,
      City.name,
      City.population,
      .....
      )
```

SQL join

Example (cont.)

```
SELECT co.name, ci.name, r.name, r.population
FROM Country co JOIN Province r
      ON co.code = r.country
      JOIN City ci
      ON r.name = ci.province AND
      ci.country=co.code
WHERE ci.population >= 5000000;
```

Tables "relevant" for query:

1. Those containing column used in projection or selection
Example: **Country**, **City**
2. Those needed to link tables of type 1.
Example: **Province**

Be careful with...

.. projection.

SQL> select * from r;

A	B	C
1	2	3
2	2	4
3	3	4

SQL> select * from s;

B	D
2	4
3	3

SQL> select * from t;

D	E
4	1
4	2

SQL> SELECT r.a, s.b FROM r JOIN s ON r.c != s.b NATURAL JOIN t;

A	B
1	2
2	2
3	2
1	2
2	2
3	2

Correct result.

BUG

QL> SELECT r.a, s.b FROM r JOIN s ON r.c != s.b NATURAL JOIN t;

A	B
1	2
1	2
2	2
2	2
3	2
3	2
2	3
2	3
3	3
3	3

Elimination of column D
in table s results in **cross
join** instead of
natural join.

wrong result!

10 Zeilen ausgewählt.

SQL / DML: joins

enhanced
SQL:1999

Natural inner join

`<tableName> NATURAL [INNER] JOIN <tableName>`

You should know what you are doing...

Example:

```
SELECT name, name
FROM City c NATURAL INNER JOIN Province r;
```

City (NAME, COUNTRY, PROVINCE, POPULATION...)
Province (NAME, COUNTRY, POPULATION...)

Result?

SQL / DML: Simple queries with joins

Inner equi-join with attribute list

`<tableName> [INNER] JOIN <tableName>
USING <attributList>`

Subset of common attributes

Example:

```
SELECT City.name, Province.name
FROM City INNER JOIN Province USING (country)
ORDER BY City.name;
```

Strange (wrong!) result...

Explain output!

Symmetry of joins or tables

Problems with **symmetric relationships**

Example: `Neighbor_of (Country1, Country2)`

Wanted: list of neighbors of Germany.

```
SELECT Country1, Country2
FROM Neighbor_of
WHERE Country1 = 'D' or Country2 = 'D'
```

Solves the problem...?

COUNTRY1	COUNTRY2
..	..
CS	D
D	NL
...	...

CASE

```
SELECT 'D' ,
(CASE
  WHEN n1.country1='D' THEN n1.country2
  WHEN n1.country2='D' THEN n1.country1
  -- Else for catch all, not needed here
END) AS "Neighbor-Country"
FROM Neighbor_of n1
WHERE n1.country1='D' OR n1.country2='D'
```

Simple solution in this case: set operator

```
(SELECT country2 as benachbart
FROM Borders
WHERE Country1='D')
UNION
(SELECT country1 as benachbart
FROM Borders
WHERE Country2='D')
```

CASE

Case expression in "target list" very useful.

```
SELECT Country.Name ,
       GNP/population AS "GNP ($) per
       Person"
```

... throws exception if `country.population = 0` (but not: `NULL`)

Can be avoided with CASE .

Outer Join

Natural outer join

`<tableName> LEFT|RIGHT|FULL
NATURAL [OUTER] JOIN <tableName>`

Outer join with condition

`<tableName> LEFT|RIGHT|FULL [OUTER] JOIN <tableName>
ON <condition>`

Example:

```
SELECT r.name, c.name
FROM Province r LEFT OUTER JOIN City c
ON (r.name = c.province AND r.C_ID = c.C_ID)
ORDER BY r.name
```

will find and output also `Provinces` without `cities`

Simple SQL: Output

Formatting the output

Different format, even HTML or other markup can be generated in some systems

"Find title, DVD_id and format for all movies"

```
BREAK ON name
COLUMN name HEADING "Land" FORMAT A15
COLUMN capital HEADING "Hauptstadt"
SELECT c.name, r.name
FROM Country c JOIN encompasses e
ON c.C_ID = e.country
JOIN Province r USING (C_ID)
WHERE e.continent LIKE 'Europ%'
ORDER BY c.name ASC;
```

Don't repeat identical titles
Column formatting
Aliases for columns

System dependent
This kind holds for
Oracle/SQL+

© HS-2010

07-DBS-SQL-25

7.2 Advanced SQL

Subselects and correlated subqueries

Using result relations instead of constants

```
SELECT name, country
FROM City
WHERE country IN ('DDR', 'D', 'BRD');
```

```
SELECT name, province
FROM City
WHERE province IN
(SELECT name
FROM Province
WHERE population > 5000000);
```

Constant list

Query dependent
constants

Independent subqueries

Independent outer
and inner SQL
block

© HS-2010

07-DBS-SQL-26

Advanced SQL: Subselects

Subqueries

Find name of country, the capital of which has less inhabitants than the capital of France.

```
SELECT name, code
FROM Country c JOIN City ci
ON c.capital=ci.name
AND c.province = ci.province
WHERE ci.population < ALL SOME [ANY]
(SELECT population
FROM City
WHERE name='Paris')
```

Wrong, if more than one result of subquery.
Needed: **value quantifier** – compare with **all** or is there **any**?

© HS-2010

07-DBS-SQL-27

Correlated Subselects

Correlated Subqueries: block structure, variables accessed in subordinate nested block

Find country name, for which the capital is at the same time the capital of a Province.

```
SELECT c.name, c.code, c.capital, c.province
FROM Country c
WHERE EXISTS
(SELECT *
FROM Province r
WHERE r.country = c.code AND r.name = c.province
AND r.capital = c.capital)
ORDER by c.name
```

Subqueries can be avoided in most cases.
How in the example above?

© HS-2010

07-DBS-SQL-28

Advanced SQL: EXISTS

NOT EXISTS extends the language

Find countries, the capital of which has a higher population than all its Provinces (except the Province of the capital)

Could be expressed in **predicate logic**:
Find country etc. such that **for all** Provinces (except that of the capital) the population is less than the population of the city.

Equivalent:

Find country etc such that **not exists** provinces different from the capital's province the population of which is equal or larger than the population of the countries capital.

© HS-2010

07-DBS-SQL-29

Advanced SQL: EXISTS

```
SELECT c.name, c.code, c.capital, c.province,
ci.population
FROM Country c JOIN City ci ON ci.name=c.capital
AND c.province = ci.name AND c.code=ci.country
WHERE NOT EXISTS
(SELECT *
FROM Province r
WHERE r.country = c.code AND c.province != r.name
AND r.population >= ci.population )
ORDER by c.name
```

© HS-2010

07-DBS-SQL-30

Division and EXISTS

Find countries which are members of **all** organizations, Germany is in.

Algebra expression?

```
SELECT DISTINCT country
FROM IsMember m1
WHERE NOT EXISTS (
  SELECT * FROM IsMember m2
  WHERE country = 'D' AND NOT EXISTS (
    SELECT * FROM IsMember m3
    WHERE m3.country = m1.country AND
          m3.organization = m2.organization
  )
)
```

© HS-2010

07-DBS-SQL-31

Set operators

Find countries which belong to Europe and Asia.

... where continent = 'Europe' ~~and~~ continent = 'Asia'

```
SELECT country FROM encompasses
WHERE continent = 'Europe'
INTERSECT
SELECT country FROM encompasses
WHERE continent = 'Asia'
```

Set operators eliminate duplicates!

<set op> ALL does not.

© HS-2010

07-DBS-SQL-32

Quantification and set operators

Example from above...

```
SELECT country      -- not in Oracle
FROM IsMember m
WHERE NOT EXISTS (
  (SELECT organization FROM IsMember
   WHERE country='D'
  )
  EXCEPT
  (SELECT organization FROM IsMember m1
   WHERE m.country = m1.country
  )
)
```

© HS-2010

07-DBS-SQL-33

Table expressions

```
SELECT name FROM country c
WHERE c.code IN (
  SELECT country FROM encompasses
  WHERE continent = 'Europe'
  INTERSECT
  SELECT country FROM encompasses
  WHERE continent = 'Asia'
)
```

Avoid subqueries, even if not correlated

© HS-2010

07-DBS-SQL-34

Table expressions

```
SELECT name FROM country c,
( SELECT country FROM encompasses
  WHERE continent = 'Europe'
  INTERSECT
  SELECT country FROM encompasses
  WHERE continent = 'Asia'
) euroAsia
WHERE euroAsia.country = c.code
```

table expression
in from list

© HS-2010

07-DBS-SQL-35

Aggregation and Grouping

Aggregate (or set) functions

f_A : table \rightarrow value, where A is some Subset of $\Sigma(\text{table})$

Aggregate functions are table functions, i.e. defined on tables or subsets of tables, not single rows

COUNT, SUM, AVG, VARIANCE, MIN, MAX are standard functions sometimes also statistical functions (e.g. variance)

© HS-2010

07-DBS-SQL-36

Using aggregation

```
SELECT to_Char( AVG (population), '9999999999.99'),
       to_char(VARIANCE(population), '99999999.99'),
       MAX(population)
FROM country
```

Target list: **do not mix aggregation and attribute values:**

```
SELECT name, MAX(population)
FROM country
```

Syntax error:
00937. 00000 - "not a single-group group function"

Extended SELECT list

Aggregate functions allow for SELECT-Blocks in target list
- provided *one* result value is guaranteed ("scalar expression").

"Total population of provinces per country":

```
SELECT code, name, (SELECT sum(population) AS EWZ
                    FROM Province r
                    WHERE r.country=c.code)
FROM Country c
```

Correlated - not a surprise....

Aggregation

How many provinces has China: easy

```
SELECT COUNT(*)
AS "Number of Provinces"
FROM Province r
WHERE c.code = 'D'
```

Table with number of provinces per country: no way

?

...	
IND	32
IL	5
CL	1
D	16
...	

Grouping of a table

...	
USA	MI
USA	MN
USA	MO
Venda	V
Vanuatu	VAN
Vietnam	VN
Volksrepublik_China	SHA
Volksrepublik_China	SHG
Volksrepublik_China	SHX
...	

... **GROUP BY** country.name

Grouping

Def.: Grouping partitions a table into groups
(or subtables) in such a way, that each group has
equal values column wise in all **grouping attributes**
GROUP BY <attr1>, ... <attrn>

The **result list** of a grouped table **may only contain**
grouping attributes or aggregations over other
attribute of the groups.

Grouping

```
SELECT c.name, c.code, COUNT(*) AS noProvinces
FROM Country c JOIN Province r
ON c.code=r.country
GROUP BY c.name, c.code
```

```
USA          USA  51
Vanuatu      VAN  1
Venda        V    1
Venezuela    YV   1
Vietnam      VN   1
Volksrepublik_China VRC 27
```

How to find the country with
- the maximum number of Provinces ?
- or those with more than one Province?

Having

Group based selection:

```

SELECT c.name, c.code, COUNT(*) AS noRegs
FROM Country c JOIN Province r ON c.code=r.country
GROUP BY c.name, c.code
HAVING COUNT(*) > 1
  
```

© HS-2010 07-DBS-SQL-43

Having

Predicates in having clause:
defined only on grouped columns or aggregated by a set function

```

SELECT c.name, c.code, SUM (ci.population) AS
"Stadtbevoelkerung"
FROM Country c JOIN City ci ON c.province =
ci.province AND c.code=ci.country
GROUP BY c.name, c.code
HAVING SUM(ci.population) > 0
  
```

- "No HAVING without 'GROUP BY' "
- Attributes in Target list must be group attributes

© HS-2010 07-DBS-SQL-44

Row and table predicates

Row predicates: evaluated for each individual row
Table predicates: evaluated on tables or groups.
 "A group / table is qualified or not"

No aggregation in row predicates: MAX, COUNT() etc do not make any sense.

Aggregation mandatory for table predicates:
 COUNT(*) > 2, MAX(population)

Brain teaser: can table predicates be expressed by row predicates?

© HS-2010 07-DBS-SQL-45

Using group by

For readability, we introduce a **VIEW**:

```

CREATE VIEW NoReg AS (
SELECT c.name, c.code, COUNT(*) AS nofRegs
FROM Country c JOIN Province r ON c.code=r.country
GROUP BY c.name, c.code )
  
```

Standard step: **construct a joined table with all the information needed.** Join previous expression with NoReg

```

SELECT c.name, c.code, c.capital, c.province, ci.population
FROM Country c JOIN City ci ON ci.name=c.capital
AND c.province = ci.province AND c.code=ci.country
JOIN NoReg n on n.name = c.name AND n.code = c.code
WHERE n.nofRegs > 1 -- now a row predicate!
AND ... -- WHERE EXISTS..
  
```

© HS-2010 07-DBS-SQL-46

Expl. cont.

```

SELECT c.name, c.code, c.capital, c.province, ci.population
FROM Country c JOIN City ci ON ci.name=c.capital
AND c.code = ci.province AND c.code=ci.country
JOIN NoReg n on n.name = c.name AND n.C_ID = c.C_ID
WHERE n.nofRegs > 1 -- now a row predicate!
AND NOT EXISTS
(
  SELECT *
  FROM Province r
  WHERE r.country = c.code AND c.province != r.province
  AND r.population > ci.population )
ORDER by c.name
  
```

NAME	C_ID	CAPITAL	R_ID	POPULATION
Daenemark	DK	Kopenhagen	DK	1358540
Grossbritannien	GB	London	ENG	6754500
Volksrepublik_China	VRC	Peking	PEK	9900000

11 rows selected

© HS-2010 07-DBS-SQL-47

Advanced SQL

Quantifiers and **counting** (in finite sets)

```

select x
from R
where EXISTS
(select * from S...)
      
```

=

```

select x
from R
where 0 <
(select count(*) from S...)
      
```

```

SELECT DISTINCT country
FROM IsMember m1
WHERE 0 =(
  SELECT Count(*) FROM IsMember m2
  WHERE country = 'D' AND NOT EXISTS (
    SELECT * FROM IsMember m3
    WHERE m3.country = m1.country AND
          m3.organization = m2.organization
  )
)
  
```

© HS-2010 07-DBS-SQL-48

GROUP BY

A realistic example¹

```
product (product_id, name, price, cost)
sales (product_id, units, date, ...)
```

"Find for each product the profit made within the last 4 weeks if less than 500 \$"

```
SELECT p.product_id, p.name,
       (sum(s.units) * (p.price - p.cost)) AS profit
FROM products p LEFT JOIN sales s USING (product_id)
WHERE s.date > CURRENT_DATE - INTERVAL '4 weeks'
GROUP BY product_id, p.name, p.price, p.cost
HAVING sum(p.price * s.units) < 500;
```

¹ from the Postgres manual

© HS-2010

07-DBS-SQL-49

SQL / DML Structuring

Temporary tables

When **inconsistency** is not an issue, temporary tables make sense

No assignment in SQL – applicative language

Instead: **Declare temporary relation** :

```
create temporary table myTmp ( ....)
```

and **assign a value** by means of an **INSERT** statement:

useful variant of **insert** → **insert** (select x,y,z from R where...)

Temporary tables are **local snapshots**, they are "dropped" at the end of a session.

© HS-2010

07-DBS-SQL-50

Structuring

Consistency thread in multiuser mode!

Structuring

Subquery factoring / local definition

```
WITH r AS (
  select m.title, m.m_id AS x, tt.m_id
  from movie m, DVD tt
  where m.year > to_date(2000, 'YYYY') )
SELECT DISTINCT r.title, t.DVD_id
from r, DVD t
where r.x = t.m_id;
```

Local definition

Compare **let** in Haskell


© HS-2010

07-DBS-SQL-51

Transitive closure

Representing a directed Graph

Example: **course prerequisites**



Represent graph by a set of nodes and a set of edges

© HS-2010

07-DBS-SQL-52

Transitive closure

Example: Find courses required for SWT

enhanced
SQL:1999

```
-- Nodes
CREATE TABLE Course(
  lnr int primary key,
  name varchar(20));
```

```
-- Edges
CREATE TABLE Requires(
  pre int references course(lnr),
  suc int references course(lnr),
  constraint req_pk primary key(pre, suc));
```

© HS-2010

07-DBS-SQL-53

ANSI SQL: Transitive closure

ANSI SQL (SQL 99) syntax for recursive traversals

```
WITH RECURSIVE PreCourse( pre, suc )
AS (SELECT pre,suc FROM Requires r WHERE pre
     NOT IN (SELECT suc FROM Requires r1)
UNION
  SELECT pre,suc
  FROM Requires r, PreCourses p
  WHERE p.suc = r.pre
)
SELECT p1.suc, c.name
FROM PreCourse p1, course c
WHERE p1.suc = c.lnr;
```

© HS-2010

07-DBS-SQL-54

Recursive processing

Querying a table recursively

- (1) Construct the table which is recursively defined
Example: **PreCourses** (**pre**, **suc**) which is the transitive closure of the **Requires** table
- (1a) Start with the "base" relation **Requires**
Requires → **PreCourses**⁰
- (1b) construct **PreCourses**ⁿ⁺¹ :
PreCoursesⁿ union [all]
additional transitive dependent tuples using
Requires and **PreCourses**ⁿ
- (2) Use constructed table (**PreCourses**) for querying

Termination?

The (iterative!) algorithm which constructs the transitive closure, terminates, if there are no new tuples to be added:

$$\text{PreCourses}^n = \text{PreCourses}^{n+1}$$

Crucial: The **result set** of the query defining the "delta" must **eventually be empty!**

In the example:

```
SELECT pre,suc
  FROM Requires r, PreCourses p
 WHERE p.suc = r.pre
```

Oracle: different mechanism

Transitive Closure in Oracle: CONNECT


Courses

LNR	NAME
1	ALP 1
2	ALP 2
3	ALP 3
4	ALP 4
5	RS
6	RO
7	Theory
8	SWP
9	DBS
10	DBProj
11	SWT

```
SELECT l.name
  FROM Course l, requires r
 WHERE r.suc = l.lnr
    START WITH pre = 1
    CONNECT BY PRIOR suc = pre;
```

```
NAME
-----
ALP 2
ALP 3
ALP 4
SWP
DBS
SWT
DBProj
```

1. Only ONE path ("start with..")
2. from leaf to root: exchange
suc and **pre**



SQL / DML: Overview

Query data

- Interactively
- Embedded in host language
important in applications, next chapter

Insert, update, delete data

7.3 SQL / DML: Update operations

Delete, Insert, Update

The easiest way to ruin your company:

```
DELETE FROM Customer;
```

deletes all rows from **Customer** relation

In general, the rows to be deleted are specified by a (search) predicate:

```
DELETE FROM <tablename> WHERE <predicate>;
```

Very different from **deleting metadata**:

```
DROP TABLE Customers; DROP SCHEMA my_database;
```

SQL / DML: Update

General form (simplified) for changing values :

```
UPDATE <tableName> SET
  <attr> = <value> [, <attr> = <value> ]0..*
 WHERE <searchPredicate>
```

- Note: without a predicate all rows will be changed

- **Primary key predicate for update** is very common

```
Update Customer SET email = 'me@acm.org'
 WHERE mem_No = 47.1; /* primary key */
Update Rental SET until_date = SYSDATE
 WHERE bike_id = 23-7789 AND c_id = 3315
 AND until_date IS NULL;
```

SQL / DML: Insertion

```
INSERT INTO <tableName> VALUES
[<value> [,value]] -- for each attribute
```

```
INSERT INTO Customer
VALUES (011, 'Müller', 'Tina', NULL, NULL, NULL);
```

Insert (2)

Incomplete form with attribute and value list:

- Order of attributes / values independent from schema
- Attributes not in value list get value NULL

```
INSERT INTO <tableName>
(attribute [,attribute]0..n) VALUES
(<value> [,value]0..n)
```

```
INSERT INTO Customer
(name, mem_no) VALUES ('Müller', 012);
```

SQL / DML: Insert data

Insertion using a query

```
INSERT INTO Foo (select * FROM Tmp)
```

Result set of query must have same type signature as table inserted to.

Bulk insertion

large file of INSERT statements may be inefficient

insertion of large data sets by specific DB tools

Postgres: COPY command to and from files (e.g. cvs)

Oracle and others: bulk loader

not standardized

SQL / DML: bulk load

Bulk load: inserting many data

Example:

```
CREATE TABLE loadtest (
  name varchar(20),
  num number(10,2));
```

loadtest.dat	
'XYZ'	4
'YZX'	5
'ZXY'	6

Oracle Syntax:

```
sqlldr <user>/<password> <controlfile>
<logfile> <badfile> <datafile>
```

loadtest.ctl

```
load data
infile 'loadtest.dat'
badfile 'loadtest.bad'
discardfile 'loadtest.dis'
APPEND INTO table loadtest
fields terminated by " , "
optionally enclosed by " ' "
(name char, num integer external)
```

Summary

- SQL: **THE interlingua** of data management
- Differences (standard, systems) considerable
- Eventually convergence towards SQL 3
- Set manipulation as dominating operation
- Set specification in a declarative way
- Grouping: frequent operation
- Many language enhancements in SQL 3 (transitive closure, structuring)
- Interactive language: embedding into host language to be discussed